

A Single-Loop Accelerated Extra-Gradient Difference Algorithm with Improved Complexity Bounds for Constrained Minimax Optimization

Yuanyuan Liu¹, Fanhua Shang², Weixin An¹, Junhao Liu¹, Hongying Liu² and Zhouchen Lin³

¹Xidian University

²Tianjin University

³Peking University



Optimization Problems

■ Constrained Minimax Optimization Problem

$$\min_{x \in \mathcal{X}} \max_{y \in \mathcal{Y}} f(x, y),$$

where $\mathcal{X} \subseteq \mathbb{R}^m$ and $\mathcal{Y} \subseteq \mathbb{R}^n$ are nonempty closed and convex feasible sets, and $f : \mathbb{R}^m \times \mathbb{R}^n \rightarrow \mathbb{R}$ is a smooth function. In recent years, this problem has drawn considerable interest from machine learning and other engineering communities such as:

- Generative adversarial networks
- Adversarial machine learning
- Game theory
- Reinforcement learning
- Empirical risk minimization
- Robust optimization

Motivations

- For NC-C minimax problems, can we design a single-loop directly accelerated algorithm with the gradient complexity lower than the best-known result $\tilde{O}(\epsilon^{-2.5})$?
- Can we propose a single-loop directly accelerated algorithm with the complexity lower than the best-known result $O(\epsilon^{-4})$ for constrained C-NC and NC-NC minimax problems?

Contributions

- We design a new single-loop accelerating algorithm for solving constrained NC-NC problems. In the proposed algorithm, we design a new extra-gradient difference scheme, and combine the gradient ascent and momentum acceleration steps for the dual variable update.
- We analyze the convergence properties of the proposed algorithm for constrained NC-NC problems. Theorem 1 shows that to find an ϵ -stationary point of f , the proposed algorithm can obtain the gradient complexity $O(\epsilon^{-2})$, which is the first time to attain the complexity bound in constrained NC-NC setting, as shown in Table 1.

Table 1: Comparison of complexities of the minimax algorithms to find an ϵ -stationary point of $f(\cdot, \cdot)$ in the NC-C, C-NC and NC-NC settings. Note that Smoothed-GDA [4] can find an ϵ -stationary point for a special problem of minimax problems with $\mathcal{O}(\epsilon^{-2})$, and $\tilde{\mathcal{O}}$ hides logarithmic factors compared to $\mathcal{O}(\cdot)$.

Optimality Criteria	References	NC-C	C-NC	NC-NC	Simplicity
Stationarity of f with smoothness and compact sets assumptions	[1]	$\tilde{\mathcal{O}}(\epsilon^{-4})$	-	-	Single-Loop
	[2]	$\tilde{\mathcal{O}}(\epsilon^{-3.5})$	-	-	Multi-Loop
	[3]	$\tilde{\mathcal{O}}(\epsilon^{-2.5})$	-	-	Multi-Loop
	[4]	$\mathcal{O}(\epsilon^{-4})$	-	-	Single-Loop
	[5]	$\mathcal{O}(\epsilon^{-4})$	$\mathcal{O}(\epsilon^{-4})$	-	Single-Loop
	This work (Theorem 1)	$\mathcal{O}(\epsilon^{-2})$	$\mathcal{O}(\epsilon^{-2})$	$\mathcal{O}(\epsilon^{-2})$	Single-Loop

- [1]** S. Lu, I. Tsaknakis, M. Hong, and Y. Chen. Hybrid block successive approximation for one-sided nonconvex min-max problems: algorithms and applications. *IEEE Trans. Signal Process.*, 68:3676–3691, 2020.
- [2]** M. Nouiehed, M. Sanjabi, T. Huang, J. D. Lee, and M. Razaviyayn. Solving a class of non-convex min-max games using iterative first order methods. In *NeurIPS*, pages 14905–14916, 2019.
- [3]** T. Lin, C. Jin, and M. I. Jordan. Near-optimal algorithms for minimax optimization. In *The Annual Conference on Learning Theory (COLT)*, pages 2738–2779, 2020.
- [4]** J. Zhang, P. Xiao, R. Sun, and Z. Luo. A single-loop smoothed gradient descent-ascent algorithm for nonconvex-concave min-max problems. In *The Annual Conference on Neural Information Processing Systems (NeurIPS)*, pages 7377–7389, 2020.
- [5]** Z. Xu, H. Zhang, Y. Xu, and G. Lan. A unified single-loop alternating gradient projection algorithm for nonconvex-concave and convex-nonconcave minimax problems. *Mathematical Programming*, 2023.

Contributions

- We also provide the theoretical analysis of our algorithm in terms of another convergence criteria (i.e., an ϵ -stationary point of φ) for constrained NC-C problems. The result shows that our algorithm improves the best-known result from $\tilde{O}(\epsilon^{-3})$ to $O(\epsilon^{-2})$, as shown in Table 2.

Table 2: Comparison of complexities of existing minimax algorithms and the proposed algorithm to find an $\phi(\cdot) := \max_{y \in \mathcal{Y}} f(\cdot, y)$ in the nonconvex-concave (NC-C) setting. This table only highlights the dependence of ϵ , and compared with $\mathcal{O}(\cdot)$, $\tilde{O}(\cdot)$ hides logarithmic factors.

NC-C Settings	References	Compact set	Complexity	Simplicity
NC-C (Stationarity of ϕ)	[1],[2]	\mathcal{X}, \mathcal{Y}	$\tilde{O}(\epsilon^{-6})$	Multi-Loop
	[3]	\mathcal{X}, \mathcal{Y}	$\tilde{O}(\epsilon^{-6})$	Single-Loop
	[4]	\mathcal{X}, \mathcal{Y}	$\tilde{O}(\epsilon^{-3})$	Multi-Loop
	[5], [6]	\mathcal{X}, \mathcal{Y}	$\tilde{O}(\epsilon^{-3})$	Multi-Loop
	This work (Theorem 2)	\mathcal{Y}	$O(\epsilon^{-2})$	Single-Loop


- [1] H. Rafique, M. Liu, Q. Lin, and T. Yang. Weakly-convex concave min-max optimization: Provable algorithms and applications in machine learning. *Optim. Method Softw.*, 2022.
- [2] C. Jin, P. Netrapalli, and M. I. Jordan. What is local optimality in nonconvex-nonconcave minimax optimization? In *International Conference on Machine Learning (ICML)*, pages 4880–4889, 2020.
- [3] T. Lin, C. Jin, and M. I. Jordan. On gradient descent ascent for nonconvex-concave minimax problems. In *International Conference on Machine Learning (ICML)*, pages 6083–6093, 2020.
- [4] K. K. Thekumprampil, P. Jain, P. Netrapalli, and S. Oh. Efficient algorithms for smooth minimax optimization. In *The Annual Conference on Neural Information Processing Systems (NeurIPS)*, pages 12659–12670, 2019.
- [5] R. Zhao. A primal dual smoothing framework for max-structured nonconvex optimization. *Mathematics of Operations Research*, 2023.
- [6] T. Lin, C. Jin, and M. I. Jordan. Near-optimal algorithms for minimax optimization. In *The Annual Conference on Learning Theory (COLT)*, pages 2738–2779, 2020.

Main Ideas of Our Method

- Gradient descent:

$$x_{t+1} = \arg \min_{x \in \mathcal{X}} \{ \langle \nabla_x f(x_t, y_t), x \rangle + \|x - x_t\|^2 / \eta_x \}. \quad (1)$$

- Extra-gradient difference prediction:

Prediction point 
$$u_{t+1/2} = y_t + \beta [2 \nabla_y f(x_t, u_{t-1/2}) - \nabla_y f(x_t, y_{t-1})]. \quad (2)$$

- Gradient ascent correction:

$$u_{t+1} = \arg \max_{u \in \mathcal{Y}} \{ \langle \nabla_y f(x_t, u_{t-1/2}), u \rangle - \|u - y_t\|^2 / \eta_y^t \}. \quad (3)$$

- Momentum acceleration:

$$y_{t+1} = \tau y_t + (1 - \tau) u_{t+1}. \quad (4)$$

Advantages of Our Algorithm and Comparison to Related Work

• **Prediction Point:** The monotonicity and co-coercivity properties of gradient operators play a crucial role for convergence analysis. However, the important properties do not hold for nonconvex problems. Some researchers developed this research in some special nonconvex settings, such as structured nonconvex and weakly convex, which require a weaker condition such as weakly monotone, pseudo-monotone, and MVI. However, such conditions seriously limit the application scope of Problem. To address this challenge, we design a new prediction point scheme, which can help us obtain a useful quasi-co-coercivity property. As a result, it does not require any monotone or structural assumption. Specifically, we find that we only require a weaker property in our theoretical analysis, that is, the co-coercivity is required at some special points $\{u_{t+1/2}, y_t\}$ ($\langle \nabla_y f(x_t, u_{t+1/2}) - \nabla_y f(x_t, y_t), u_{t+1/2} - y_t \rangle \geq \rho \|\nabla_y f(x_t, u_{t+1/2}) - \nabla_y f(x_t, y_t)\|^2$ with $\rho > 0$). Thus, we develop a decoupling idea to construct the prediction point $u_{t+1/2}$. That is, we use the gradients w.r.t. y at $u_{t-1/2}, y_{t-1}$ instead of those at the points $u_{t+1/2}, y_t$. We can obtain a property (called the quasi-co-coercivity), which plays a key role in our theoretical analysis.

Property 1 (Quasi-Cocoercivity) *Let $u_{t+1/2}$ be updated in Eq. (3), then*

$$\langle 2\nabla_y f(x_t, u_{t-1/2}) - \nabla_y f(x_t, y_{t-1}), u_{t+1/2} - y_t \rangle = \beta \|2\nabla_y f(x_t, u_{t-1/2}) - \nabla_y f(x_t, y_{t-1})\|^2.$$

Core Propositions

Proposition 1 (Upper bound of primal-dual updates). Suppose Assumption 1 holds. Let $\{(x_t, y_t, u_t)\}$ be a sequence generated by Algorithm 1

$$\begin{aligned} & f(x_{t+1}, u_{t+1/2}) - f(x_t, y_t) \\ & \leq - \left(\frac{1}{\eta_x} - \frac{L}{2} \right) \|x_{t+1} - x_t\|^2 + L \|u_{t+1/2} - y_t\|^2 + L \|y_t - y_{t-1}\|^2 + \underbrace{\langle \nabla_y f(x_t, y_{t-1}), u_{t+1/2} - y_t \rangle}_{A_1}. \end{aligned}$$

Proposition 2 (Upper bound of dual updates). Suppose Assumption 1 holds. Let $\{(x_t, y_t, u_t)\}$ be a sequence generated by Algorithm 1 then

$$\begin{aligned} & f(x_{t+1}, y_{t+1}) - f(x_{t+1}, u_{t+1/2}) \\ & \leq \underbrace{\tau \langle \nabla_y f(x_t, u_{t-1/2}), y_t - u_{t+1} \rangle}_{A_2} + \underbrace{\langle \nabla_y f(x_t, u_{t-1/2}), u_{t+1} - u_{t+1/2} \rangle}_{A_3} + a_t \end{aligned}$$

Convergence Guarantees

$$G_t := f(x_t, y_t) + 9L\|u_t - y_{t-1}\|^2 + 8L\beta^2 \|\nabla_y f(x_{t-1}, u_{t-3/2}) - \nabla_y f(x_{t-1}, y_{t-2})\|^2.$$

Next, we need to prove that our defined potential function can make sufficient decrease at each iteration, i.e., $G_t - G_{t+1} > 0$ as in Lemma 1 below. To prove Lemma 1, we will provide and prove the following upper bounds.

$$\begin{aligned} G_{t+1} - G_t &:= \underbrace{f(x_{t+1}, y_{t+1}) - f(x_{t+1}, u_{t+1/2})}_{\text{Proposition 1}} + \underbrace{f(x_{t+1}, u_{t+1/2}) - f(x_t, y_t)}_{\text{Proposition 2}} + \text{other terms} \\ &= \underbrace{A_1 + A_2 + A_3}_{\text{Quasi-Cocoercivity in Property 1}} + \text{other terms.} \end{aligned}$$

Lemma 1 (Descent estimate of G). *Suppose Assumption 1 holds. Let $\{(x_t, y_t, u_t)\}$ be a sequence generated by Algorithm 1 then*

$$\begin{aligned} &G_0 - G_T \\ &= \sum_{t=0}^{T-1} (G_t - G_{t+1}) \\ &\geq \sum_{t=0}^{T-1} \left(\frac{1}{2\beta} \|u_{t+1} - y_t\|^2 + \frac{\beta - 30L\beta^2}{2} \|\nabla_y f(x_t, u_{t-1/2}) - \nabla_y f(x_t, y_{t-1})\|^2 + \frac{1}{2\eta_x} \|x_{t+1} - x_t\|^2 \right). \end{aligned}$$

Theorem 1 (Stationarity of f in constrained NC-NC settings) *Suppose Assumptions 1 and 2 hold. Then the complexity of Algorithm 1 to find an ϵ -stationary point of f with $\eta_x \leq \frac{1}{4L}$, $\beta \leq \frac{1}{60L}$ and $\tau \geq 1/2$ is bounded by*

$$\mathcal{O}\left(\frac{G_0 - \underline{G} + 2LD_y^2}{\epsilon^2}\right),$$

where $G_0 := G(x_0, y_0)$, and $\underline{G} := \min_{x \in \mathcal{X}} \phi(x)$.

For constrained NC-NC problems, the gradient complexity of our EGDA algorithm to find an ϵ -stationary point of f is $\mathcal{O}(\epsilon^{-2})$. That is, our EGDA algorithm is first to obtain the gradient complexity in constrained NC-NC setting.

Theorem 2 (Stationarity of ϕ for constrained NC-C settings) *Using the same notation as in Theorem 1 and f is concave with respect to y . Let $\{(x_t, y_t, u_t)\}$ be a sequence generated by Algorithm 1. Then the gradient complexity of Algorithm 1 to find an ϵ -stationary point of ϕ is bounded by*

$$\mathcal{O}\left(\frac{D_{\mathcal{Y}}(G_0 - \underline{G} + 2LD_{\mathcal{Y}}^2)}{\epsilon^2}\right).$$

That is, the proposed algorithm can improve the best-known gradient complexity from $\tilde{\mathcal{O}}(\epsilon^{-3})$ to $\mathcal{O}(\epsilon^{-2})$.

Experimental Results

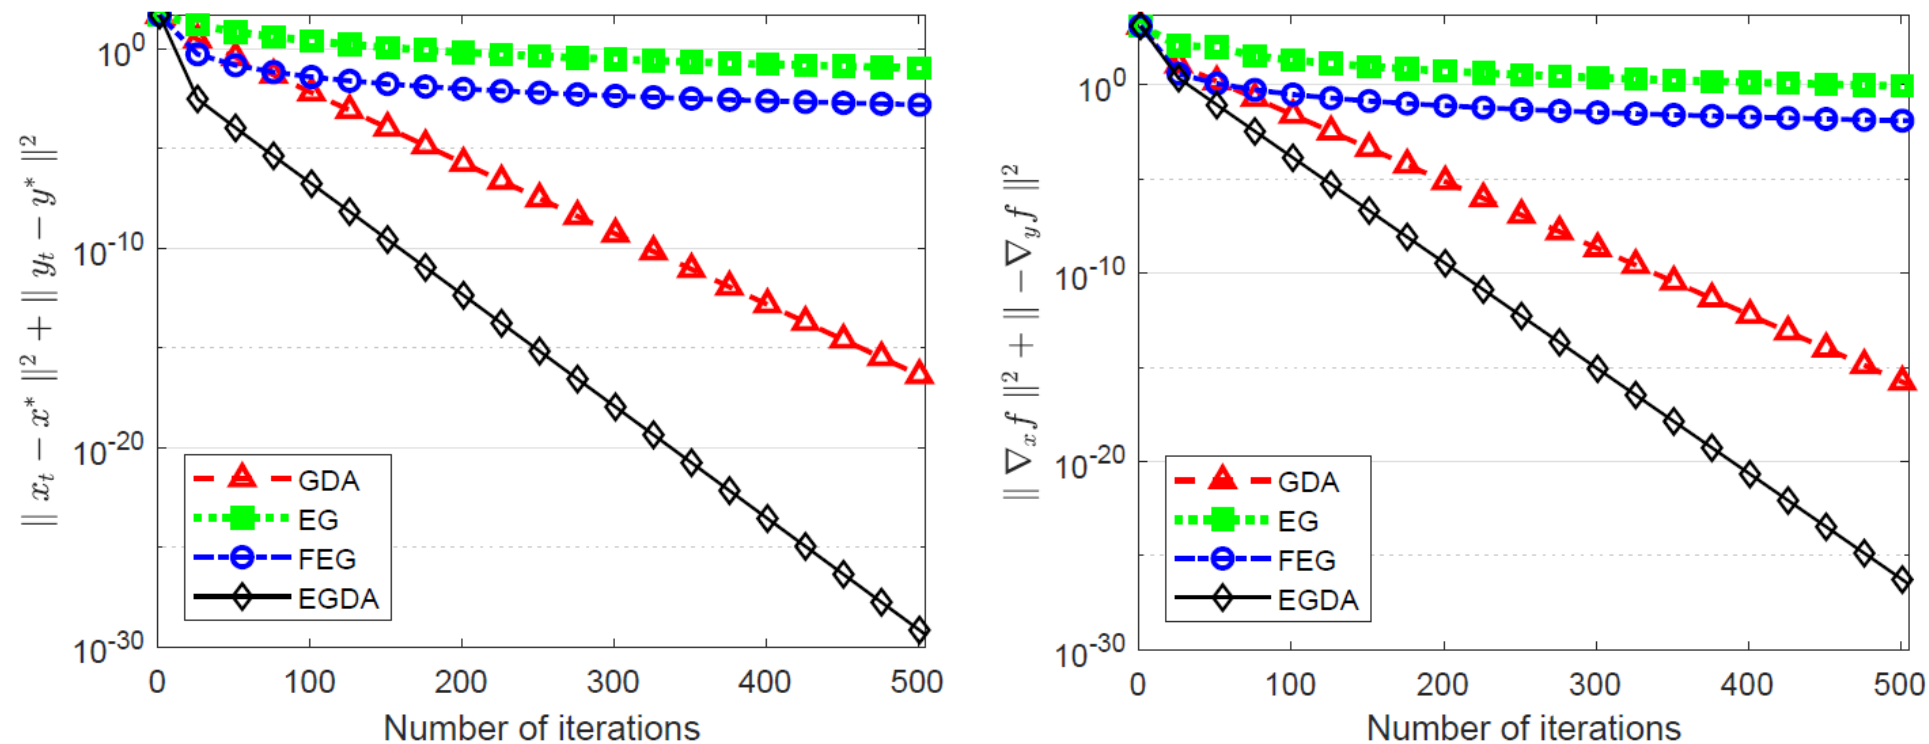


Figure 1: Comparison of all the methods for solving the NC-NC problem, $f(x, y) = x^2 + 3 \sin^2 x \sin^2 y - 4y^2 - 10 \sin^2 y$. Left: Convergence in terms of $\|x_t - x^*\|^2 + \|y_t - y^*\|^2$, where (x^*, y^*) is the global saddle point; Right: Convergence in terms of $\|\nabla_x f\|^2 + \|\nabla_y f\|^2$.


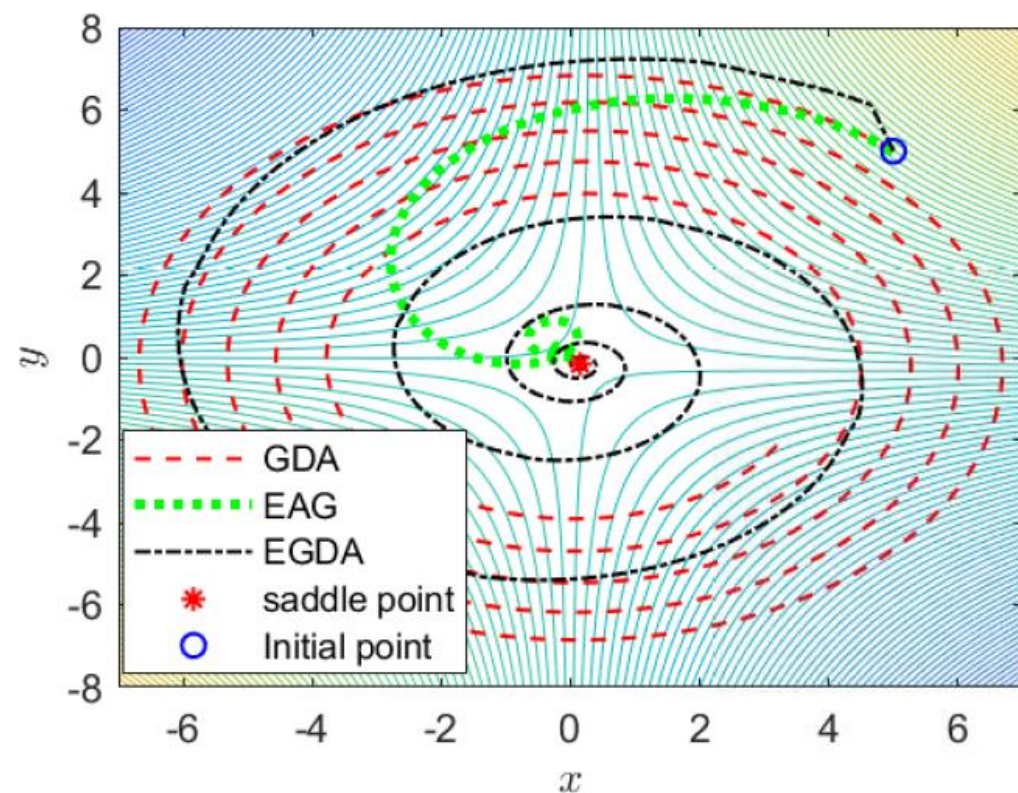


Figure 2: Comparison of all the methods for solving the convex-concave problem, $f(x, y) = \log(1 + e^x) + 3xy - \log(1 + e^y)$. Left: Trajectories of the three algorithms; Right: Convergence in terms of $\|x_t - x^*\|^2 + \|y_t - y^*\|^2$.

Experimental Results

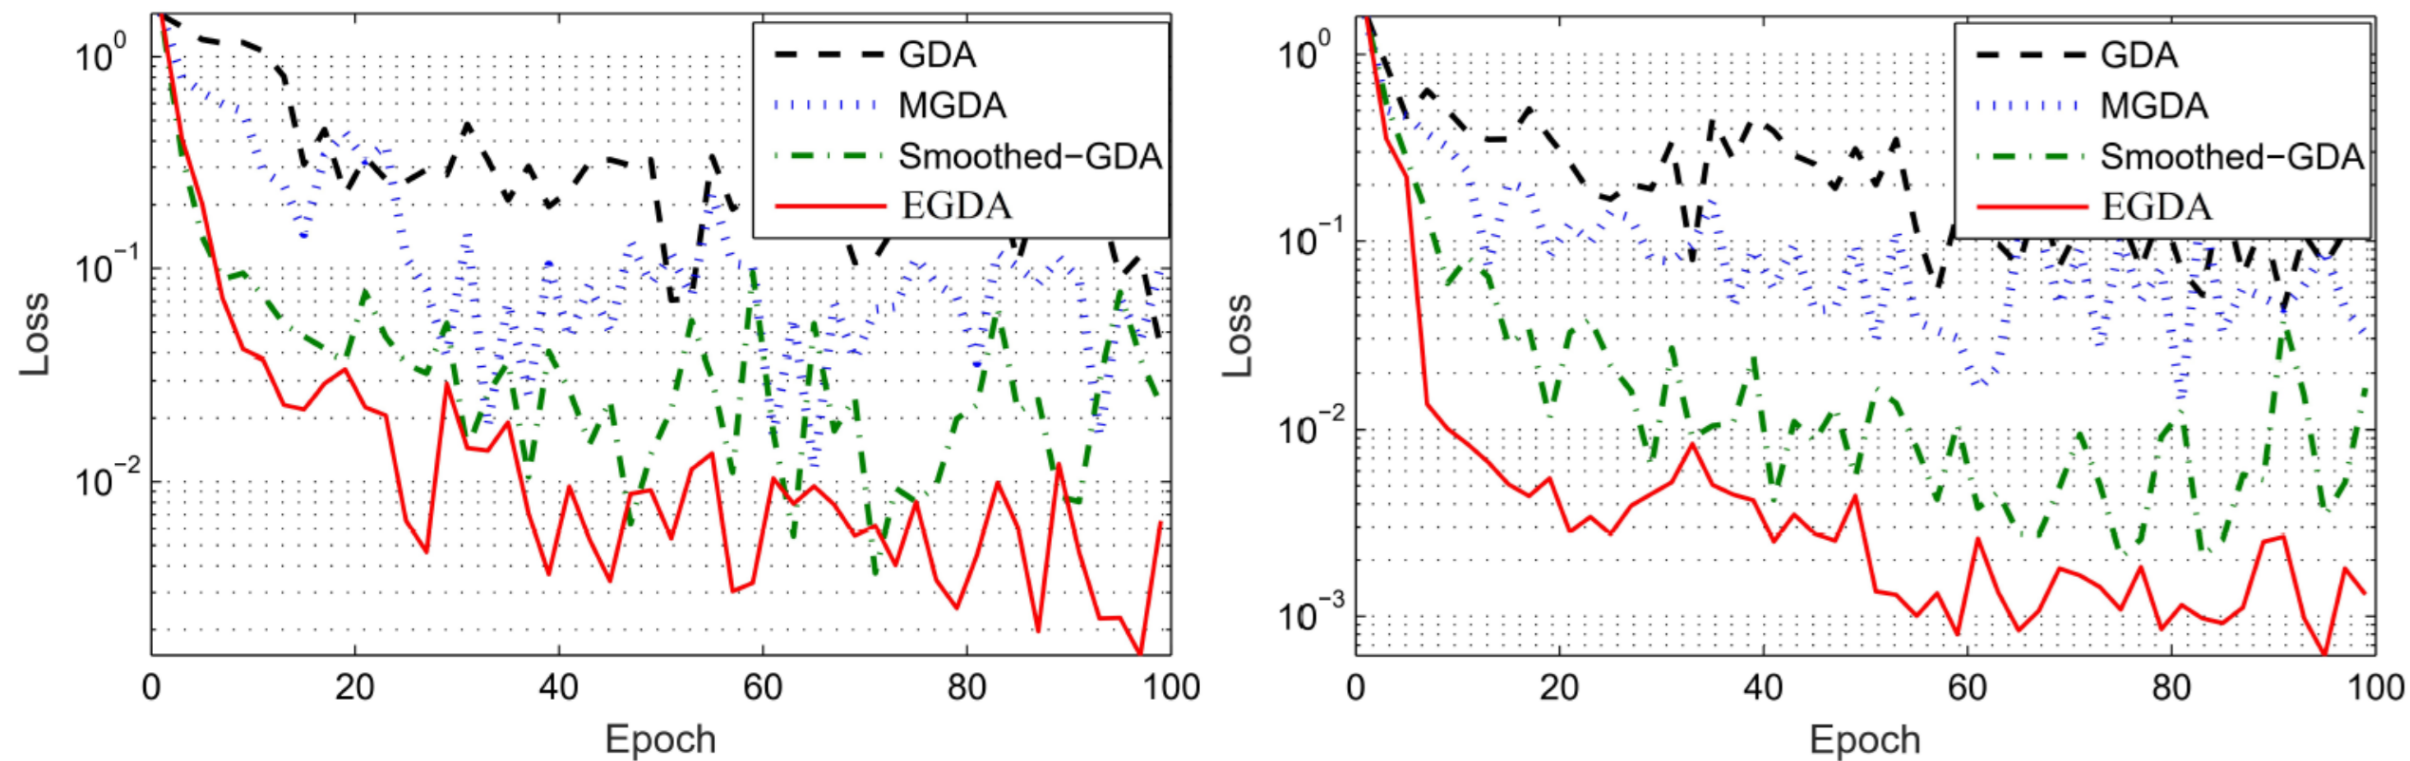


Figure 3: Convergence speed of all the algorithms on Fashion MNIST (left) and MNIST (right).

Thank you!